

rolando_lab6

February 24, 2023

1 Lab 6 - EDA with Clustering

1.0.1 Jackson Rolando

1.1 Part 1 - Load and Transform the Data

We'll load the data, split it into training and testing, transform it into a sparse bag of words matrix, excluding words appearing less than 10 times:

```
[1]: import glob
import json
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: objects = []
for file in glob.glob('./email_json/*.json'):
    with open(file) as f:
        objects.append(json.load(f))

df = pd.DataFrame(objects)
df.head()
```

```
[2]: category                                to_address \
0      ham  BREAKINGNEWS Subscribers<BREAKINGNEWS-Subscrib...
1      spam                                <theorize@plg.uwaterloo.ca>
2      spam                "Theorize" <theorize@plg.uwaterloo.ca>
3      spam                warwickktwarwic@speedy.uwaterloo.ca
4      ham                                R-help@stat.math.ethz.ch

                                from_address \
0      BREAKING NEWS<breakingnews@foxnews.com>
1      "cschai" <cschai@syhmco.co.kr>
2      "Aegis Capital Group LLC" <Estela.Burch@smapxs...
3      "shar Nobis" <sharNobis@autotradebuyer.co.uk>
4      jessica.gervais@tudor.lu

                                subject \
0      FNC Alert
```

```

1                                     rtfmub
2  Invitation to fill in the vacant position of a...
3                                     Terrific gains possible!
4                                     [R] time serie generation

```

```

                                     body
0  PELOSI, REID SIGN WAR-SPENDING BILL THAT INCLU...
1  \n\n\n\n\n\n\n\nwyat\nlnpmoqrkhapibcegd\n\n\n\n...
2  \n\n\n\nDear    sirs,\nAegis      Capital Gro...
3  http://s6.bilder-hosting.de/img/7LR4W.jpg\nImp...
4  \nDear all,\n\nI would like to generate a regu...

```

```
[3]: df.category = df.category.astype('category')
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 63542 entries, 0 to 63541
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   category        63542 non-null  category
1   to_address      63141 non-null  object
2   from_address    63542 non-null  object
3   subject         63410 non-null  object
4   body            63542 non-null  object
dtypes: category(1), object(4)
memory usage: 2.0+ MB

```

```
[4]: # from sklearn.model_selection import train_test_split

# data = df.drop(['category'], axis=1)
# labels = df['category']

# train, test, train_label, test_label = train_test_split(data, labels,
↳ test_size=.15)

# train.to_json('./data/train_data.json')
# test.to_json('./data/test_data.json')
# train_label.to_json('./data/train_label.json')
# test_label.to_json('./data/test_label.json')

train = pd.read_json('./data/train_data.json')
test = pd.read_json('./data/test_data.json')
train_label = pd.read_json('./data/train_label.json', typ='series')
test_label = pd.read_json('./data/test_label.json', typ='series')

```

```
[5]: from sklearn.feature_extraction.text import CountVectorizer
```

```
vectorizer = CountVectorizer(binary=True, min_df=10)
train_feat_mat = vectorizer.fit_transform(train.body)
```

```
test_feat_mat = vectorizer.transform(test.body)
```

```
print(train_feat_mat.shape)
print(test_feat_mat.shape)
```

```
(54010, 29326)
```

```
(9532, 29326)
```

1.2 Part 2 - Clustering the Emails

I chose KMeans, as the clusters should be linearly separable. Knowing how the data looks from the previous lab, the density of the data, as well as the mixture of ham vs spam, seems to change throughout the space, so I'll keep the number of clusters high.

We'll run SVD dimension reduction to bring the columns down to 10, then cluster the data and make a map from cluster to label for predictions:

```
[6]: from sklearn.decomposition import TruncatedSVD
```


```
condenser = TruncatedSVD(n_components=10)
train_condensed_mat = condenser.fit_transform(train_feat_mat)
```

```
test_condensed_mat = condenser.transform(test_feat_mat)
```

```
print(condenser.explained_variance_ratio_)
plt.bar(np.arange(condenser.explained_variance_ratio_.size), condenser.
    ↪ explained_variance_ratio_)
```

```
[0.05780993 0.02786145 0.01750974 0.01319702 0.01363418 0.01094181
 0.01059009 0.01014235 0.00814024 0.0079912 ]
```

```
[6]: <BarContainer object of 10 artists>
```



```
[104]: from sklearn.cluster import KMeans
```

```
clusterer = KMeans(n_clusters=20)
```

```
train_clusters = clusterer.fit_predict(train_condensed_mat[:, 0:2])
```

```
test_clusters = clusterer.predict(test_condensed_mat[:, 0:2])
```

```
/home/rolo/.local/lib/python3.8/site-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
```

```
[105]: np.unique(train_clusters)
```

```
[105]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
          17, 18, 19], dtype=int32)
```


```
[109]: colors = {0: 'blue', 1: 'orange', 2: 'green', 3: 'magenta', 4: 'cornflowerblue',
           5: 'yellow', 6: 'red', 7: 'aqua'}
```

```
fig, ax = plt.subplots()
```

```

for cluster_num in np.unique(train_clusters):
    indices = np.where(train_clusters == cluster_num)
    ax.scatter(train_condensed_mat[indices, 0], train_condensed_mat[indices, 1], label=cluster_num, c=colors[cluster_num % len(colors)])
ax.legend()
plt.show()

```



We'll check the counts of ham vs spam in each of the clusters:

```

[110]: total = train_label.shape[0]
ham_num = (train_label == 'ham').sum()
format_percent = '%.2f'%(100 * ham_num / total)
print(f'total: {total} - ham: {ham_num} - {format_percent}%\n')

cluster_hams = {}
for clust in np.unique(train_clusters):
    total = train_clusters[train_clusters == clust].shape[0]
    ham_num = train[(train_clusters == clust) & (train_label == 'ham')].shape[0]

    dec_ham = ham_num / total

```

```

cluster_hams[clust] = dec_ham

format_percent = '%.2f'%(100 * dec_ham)
print(f'cluster {clust}: {total} - ham: {ham_num} - {format_percent}%')

print(f'\n{cluster_hams}')

clust_label_map = {}
for clust, percent_ham in cluster_hams.items():
    clust_label_map[clust] = 'ham' if percent_ham > 0.5 else 'spam'

print(f'\n{clust_label_map}')

def cluster_predict(condensed_mat):
    clusters = clusterer.predict(condensed_mat)

    mapped = np.vectorize(lambda cluster_num: clust_label_map.
        ↳get(cluster_num))(clusters)

    return mapped

train_predicted = cluster_predict(train_condensed_mat[:, 0:2])
train_predicted

```

total: 54010 - ham: 19581 - 36.25%

```

cluster 0: 1857 - ham: 1857 - 100.00%
cluster 1: 5814 - ham: 986 - 16.96%
cluster 2: 5567 - ham: 826 - 14.84%
cluster 3: 706 - ham: 694 - 98.30%
cluster 4: 3673 - ham: 542 - 14.76%
cluster 5: 216 - ham: 186 - 86.11%
cluster 6: 753 - ham: 73 - 9.69%
cluster 7: 8343 - ham: 1423 - 17.06%
cluster 8: 1342 - ham: 1342 - 100.00%
cluster 9: 1670 - ham: 326 - 19.52%
cluster 10: 441 - ham: 441 - 100.00%
cluster 11: 2077 - ham: 1979 - 95.28%
cluster 12: 375 - ham: 297 - 79.20%
cluster 13: 5745 - ham: 756 - 13.16%
cluster 14: 5568 - ham: 516 - 9.27%
cluster 15: 1293 - ham: 1293 - 100.00%
cluster 16: 1993 - ham: 1840 - 92.32%
cluster 17: 2749 - ham: 605 - 22.01%
cluster 18: 1500 - ham: 1385 - 92.33%
cluster 19: 2328 - ham: 2214 - 95.10%

```

```
{0: 1.0, 1: 0.1695906432748538, 2: 0.14837434884138676, 3: 0.9830028328611898,
4: 0.14756329975496868, 5: 0.8611111111111112, 6: 0.09694555112881806, 7:
0.17056214790842622, 8: 1.0, 9: 0.19520958083832335, 10: 1.0, 11:
0.9528165623495426, 12: 0.792, 13: 0.13159268929503917, 14: 0.09267241379310345,
15: 1.0, 16: 0.9232313095835424, 17: 0.22008002910149146, 18:
0.9233333333333333, 19: 0.9510309278350515}
```

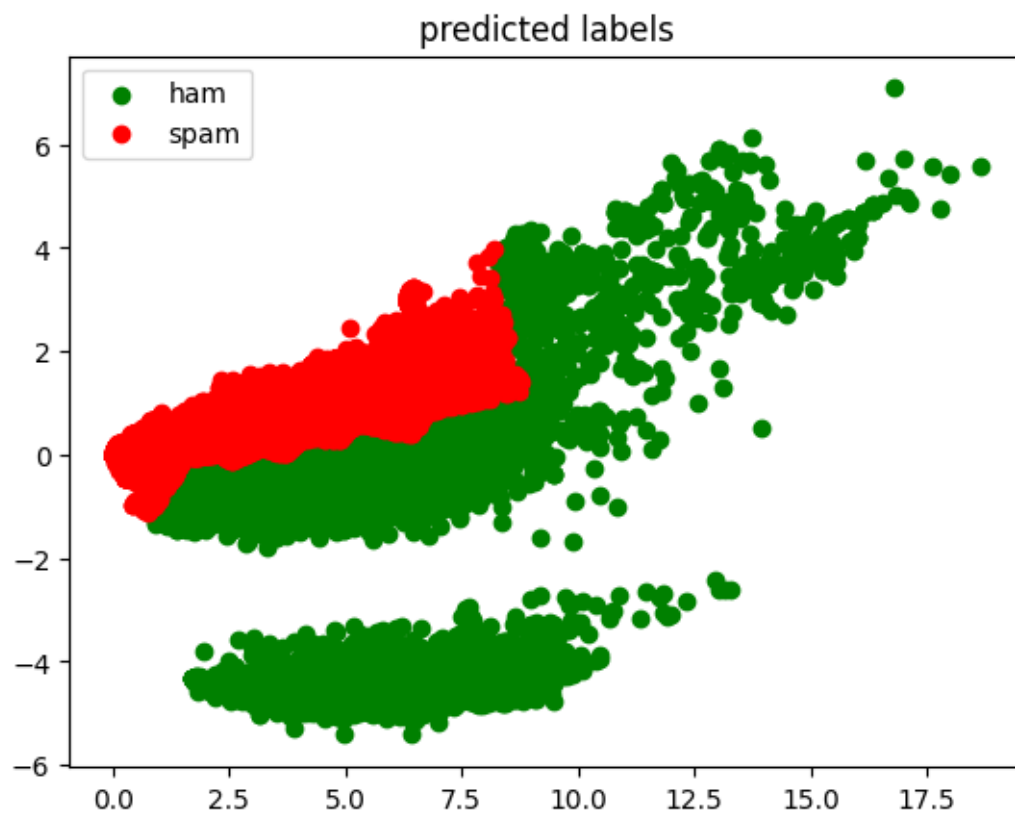
```
{0: 'ham', 1: 'spam', 2: 'spam', 3: 'ham', 4: 'spam', 5: 'ham', 6: 'spam', 7:
'spam', 8: 'ham', 9: 'spam', 10: 'ham', 11: 'ham', 12: 'ham', 13: 'spam', 14:
'spam', 15: 'ham', 16: 'ham', 17: 'spam', 18: 'ham', 19: 'ham'}
```


```
[110]: array(['spam', 'ham', 'spam', ..., 'spam', 'spam', 'ham'], dtype='<U4')
```

```
[111]: colors = {'ham': 'green', 'spam': 'red'}

fig, ax = plt.subplots()
for label in np.unique(train_label):
    indices = np.where(train_predicted == label)
    ax.scatter(train_condensed_mat[indices, 0], train_condensed_mat[indices, 1],
               label=label, c=colors[label])
ax.set_title('predicted labels')
ax.legend()
plt.show()

fig, ax = plt.subplots()
for label in np.unique(train_label):
    indices = np.where(train_label == label)
    ax.scatter(train_condensed_mat[indices, 0], train_condensed_mat[indices, 1],
               label=label, c=colors[label])
ax.set_title('actual labels')
ax.legend()
plt.show()
```





With 20 clusters, the algorithm was able to generalize the area in which most spam emails occur. Now we'll predict on the test data and make a confusion matrix based on its results.


```
[121]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

test_cluster_preds = cluster_predict(test_condensed_mat[:, 0:2])
print(test_cluster_preds.shape[0])
print((test_cluster_preds == 'ham').sum())

cm = confusion_matrix(test_label, test_cluster_preds)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['ham',
↪ 'spam'])

disp.plot()
plt.show()
```

9532
2487



```
[124]: print(f'{{(2382 + 5916)}} out of {{(2382 + 5916 + 1129 + 105)}} predicted_
        ↪correctly')
```

8298 out of 9532 predicted correctly

It looks like the clustering performed well for the test set.

1.3 Part 3 - Calculating Document Frequencies of Words

```
[129]: ham_mat = train_feat_mat[train_predicted == 'ham']
        spam_mat = train_feat_mat[train_predicted == 'spam']

        from scipy.sparse import csc_matrix
        ham_mat = csc_matrix(ham_mat)
        spam_mat = csc_matrix(spam_mat)
```

```
[134]: ham_word_cnts = ham_mat.sum(axis=0)
        print(ham_word_cnts.shape)
        spam_word_cnts = spam_mat.sum(axis=0)
        print(spam_word_cnts.shape)
```

(1, 29326)

(1, 29326)

```
[149]: def print_freqs(word):
        word_ind = vectorizer.vocabulary_[word]
        freq = ham_word_cnts[0, word_ind]
        total = ham_mat.shape[0]
        format_percent = '%.2f'%(100 * freq / total)
        print(f'{word}:\n ham: {freq} - appears in {format_percent}% of
        ↪ham-predicted documents')

        freq = spam_word_cnts[0, word_ind]
        total = spam_mat.shape[0]
        format_percent = '%.2f'%(100 * freq / total)
        print(f' spam: {freq} - appears in {format_percent}% of spam-predicted
        ↪documents\n')
```

```
[150]: print_freqs('love')
        print_freqs('works')
        print_freqs('different')
```

love:

ham: 355 - appears in 2.51% of ham-predicted documents
spam: 1365 - appears in 3.42% of spam-predicted documents

works:

ham: 1407 - appears in 9.96% of ham-predicted documents
spam: 1179 - appears in 2.96% of spam-predicted documents

different:

ham: 1497 - appears in 10.60% of ham-predicted documents
spam: 936 - appears in 2.35% of spam-predicted documents

1.4 Part 4 - Find Enriched Words with Statistical Testing

```
[222]: import scipy.stats as stats

        enriched_in_0 = np.vectorize(lambda spam_word_cnt, ham_word_cnt: stats.
        ↪binom_test(ham_word_cnt, ham_mat.shape[0], (spam_word_cnt / spam_mat.
        ↪shape[0]), alternative='greater'))(spam_word_cnts, ham_word_cnts)
        enriched_in_1 = np.vectorize(lambda spam_word_cnt, ham_word_cnt: stats.
        ↪binom_test(spam_word_cnt, spam_mat.shape[0], (ham_word_cnt / ham_mat.
        ↪shape[0]), alternative='greater'))(spam_word_cnts, ham_word_cnts)
```

```
[224]: print(enriched_in_0.shape)
        print(enriched_in_1.shape)
```

(1, 29326)

(1, 29326)

```
[228]: word = 'works'
ind = vectorizer.vocabulary_[word]
print(f'{word}: \nin 0: {enriched_in_0[0, ind]}')
print(f'in 1: {enriched_in_1[0, ind]}')
print()

word = 'love'
ind = vectorizer.vocabulary_[word]
print(f'{word}: \nin 0: {enriched_in_0[0, ind]}')
print(f'in 1: {enriched_in_1[0, ind]}')
print()
```

```
works:
in 0: 0.0
in 1: 0.9999999999999999

love:
in 0: 0.9999999997999472
in 1: 1.8293020464383167e-28
```

Works is enriched in 0, and love is enriched in 1. So, works is more often in ham, and love is more often in spam.

```
[259]: words_enriched_in_0 = []
words_enriched_in_1 = []
for word, index in vectorizer.vocabulary_.items():
    if(word.isalpha()):
        p_0 = enriched_in_0[0, index]
        p_1 = enriched_in_1[0, index]
        ham_cnt = ham_word_cnts[0, index]
        spam_cnt = spam_word_cnts[0, index]
        words_enriched_in_0.append((p_0, word, ham_cnt))
        words_enriched_in_1.append((p_1, word, spam_cnt))
```

```
[256]: print('enriched in 0:')
for word in sorted(words_enriched_in_0, key=lambda x: x[0])[0:200]:
    print(word[1])
```

```
enriched in 0:
hello
new
and
http
www
with
```

need
for
any
this
to
try
be
me
in
the
kernel
module
so
can
use
my
dectalk
linux
running
that
from
am
there
would
have
some
something
help
just
files
ve
using
working
want
could
of
same
thanks
do
mail
speakup
mailing
list
braille
uwo
speech
mailman
listinfo

which
or
not
please
does
system
also
get
it
at
if
code
is
original
message
bounces
stat
math
ethz
ch
mailto
on
sent
subject
question
hi
context
statistics
other
example
as
mean
read
but
alternative
html
version
deleted
https
posting
guide
project
org
provide
commented
minimal
self
contained

reproducible
an
when
plot
format
way
wed
may
wrote
how
library
matrix
pairwise
dataset
output
each
function
following
like
problem
file
advance
university
brian
ripley
stats
ox
professor
applied
oxford
tel
fax
written
cc
cheers
lists
ibiblio
where
make
used
possible
root
rev
work
works
learn
what
error

unsubscribe
view
source
two
user
os
data
first
trouble
pl
prostoalex
writes
looks
instead
then
method
additional
methods
moskalyuk
binary
etc
seems
isn
archive
computerworld
least
maybe
end
command
viewarticlebasic
mises
functions
argument
tid
id
begin
since
latimes
done
signed
true
seem
add
skype
think
doesn
pcworld
gam

run
type

```
[257]: print('enriched in 1:')  
       for word in sorted(words_enriched_in_1, key=lambda x: x[0])[0:200]:  
           print(word[1])
```

enriched in 1:
visit
our
online
shipping
worldwide
lorena
your
loan
refinance
credit
ready
lenders
established
lowest
payments
thay
unhappy
lover
safest
pnis
anywhere
millions
men
enhan
ement
hes
atches
deliver
product
unreal
dis
ounts
pybal
hk
his
he
producttestpanel
speedy
uwaterloo
sternshirt

buy
high
quality
price
us
photoshop
canadian
sterling
thrive
industrious
fasten
delivery
customer
day
instructed
money
gram
instruct
literacy
sociability
expenses
chargers
traveling
vacancy
sydneycentre
vacancies
glad
yours
rcs
howstuffworks
prices
suite
click
younger
low
news
instant
premiere
safe
fast
alerts
chose
discounts
meds
oem
packing
booklets
macromedia

adobe
acrobat
corel
grafix
illustrator
ableton
gigabook
borland
gervasio
her
centrecar
viagra
erection
sex
life
suffer
difficulties
maintaining
erectile
dysfunction
lotteryagent
lotto
renown
powerballmegamillionseuromillions
lotterylotto
ennis
wakes
fiv
poo
shook
choky
mornin
arched
sheep
tent
lagging
commutin
aint
wouldnt
bein
mrs
solemnly
him
bled
danglars
thoughtfully
viscount
corporal

itch
strod
dine
caderousse
furnish
whine
uttered
tasteless
wove
repulsive
beauchamp
bovine
pasha
haide
noirtier
god
grieving
she
procession
strove
dodo
nursery
throng
pills
itemyour
soft
tabs
cialis
jelly
levitra
startet
firma
preis
wkn
isin
markt
sie
rallye
ist
alert
alerting
aqui
loss
blurtd
macbain
drained
yourself
admirrante

fda
pharmacies
pharmacy
professional
soma
growth
hormone
meridia
tramadol
herbalking
herbal
penis
enlargement
tutorials
ineffective
pumps

1.5 Reflection Questions

1. There are many words in the English language that are necessary for basic sentence structure, like 'the', 'and', 'a'... that are sure to appear in most emails, regardless of if it's ham or spam. The separate cluster consists of words that were not seen in spam at all, potentially being terms directly related to specified work topics, maybe people's names, or just overall great vocabulary not usually seen in foreign-written spam emails.
2. In my clustering, there is both ham and spam in most of the clusters. I tried to achieve a tighter fitting clustering to the data. There were a few clusters which covered the bottom section of points, which only contained ham messages. The mapping could be changed to make sure that ONLY these clusters are mapped to ham.
3. The ham messages contained a lot of 'advertisey' language, trying to make whatever they're talking about seem too good to be true. The ham messages seem to be more topical, and sometimes contain typos.

```
[261]: train[train_predicted == 'ham'][0:25]
```

```
[261]:                                     to_address \
42791 "Speakup is a. screen review system for Linux...
34315 "Y G" <gatemaze@gmail.com>, <r-help@stat.math...
54248      Felicity Jones <felicity.jones@stanford.edu>
35193                                     cc-community@lists.ibiblio.org
37793                                     avcooper@flax9.uwaterloo.ca
32395                                     r-help@stat.math.ethz.ch
56651      "Christopher R. Hertel" <crh@ubiqx.mn.org>
30652                                     perl6-internals@perl.org
49556      Chabot Denis <chabotd@globetrotter.net>
434      Christoph Buser <buser@stat.math.ethz.ch>
61015 "Speakup is a screen review system for Linux."...
53029      "H. Paul Benton" <hpbenton@scripps.edu>
1066      Ben Edwards <funkytwig@gmail.com>
```

44225 James Peach <jpeach@samba.org>
 21039 r-help@stat.math.ethz.ch
 62338 msn-list@te.verweg.com
 2380 "Sundar Dorai-Raj" <sundar.dorai-raj@pdf.com>
 33977 samba-cvs@samba.org
 4286 "Speakup is a screen review system for Linux."...
 28368 ted.harding@manchester.ac.uk, r-help@stat.math...
 25997 bugs-bitbucket@netlabs.develooper.com
 13420 r-help@stat.math.ethz.ch
 44333 "Darren Duncan" <darren@darrenduncan.net>
 62197 cc-community@lists.ibiblio.org
 21450 samba-technical@lists.samba.org

from_address \

42791 Radical NetSurfer <radsurfer@yahoo.com>
 34315 "Leeds, Mark \ (IED\)" <Mark.Leeds@morganstanle...
 54248 Prof Brian Ripley <ripley@stats.ox.ac.uk>
 35193 Andy Kaplan-Myrth <techlaw@uottawa.ca>
 37793 slashdot@slashdot.org
 32395 (Ted Harding) <ted.harding@nessie.mcc.ac.uk>
 56651 Jeremy Allison <jra@samba.org>
 30652 allison@cvs.develooper.com
 49556 Prof Brian Ripley <ripley@stats.ox.ac.uk>
 434 "Daniel Tahin" <e0226781@student.tuwien.ac.at>
 61015 Gaijin <gaijin@clearwire.net>
 53029 "jim holtman" <jholtman@gmail.com>
 1066 Srinivas <srini@geekcrossing.net>
 44225 "Gerald (Jerry) Carter" <jerry@samba.org>
 21039 Leonardo Lami <lami@faunalia.it>
 62338 msn-list-request@te.verweg.com
 2380 "Deepayan Sarkar" <deepayan.sarkar@gmail.com>
 33977 abartlet@samba.org
 4286 "Littlefield, tyler" <compgeek13@gmail.com>
 28368 Stephen Tucker <brown_emu@yahoo.com>
 25997 Steve Peters (via RT) <parrotbug-followup@parr...
 13420 francogrex <francogrex@mail.com>
 44333 "Jonathan Lang" <dataweaver@gmail.com>
 62197 Terry Hancock <hancock@anansinspaceworks.com>
 21450 Jeremy Allison <jra@samba.org>

subject \

42791 Help with dtpc.o for Centos 5
 34315 Re: [R] general question about use of list
 54248 Re: [R] sizing and saving graphics in R
 35193 [cc-community] Podcasting Legal Guide for Canada
 37793 [Slashdot] Stories for 2007-05-23
 32395 Re: [R] to draw a smooth arc

56651 Re: Preserving NTFS permissions.
 30652 [svn:parrot-pdd] r18213 - trunk/docs/pdds
 49556 Re: [R] Reducing the size of pdf graphics file...
 434 Re: [R] Preconditions for a variance analysis
 61015 Re: gentoo dropping speakup support
 53029 Re: [R] data type for block data?
 1066 Re: Passing multiple mixed arguments to subs
 44225 Re: svn commit: samba r23689 - in branches/SAM...
 21039 [R] select row
 62338 MSN-list Digest, Vol 7, Issue 90
 2380 Re: [R] Positioning in xyplot
 33977 svn commit: samba r23695 - in branches/SAMBA_4...
 4286 Re: Slackware 11 aliases, anyone?
 28368 Re: [R] Tools For Preparing Data For Analysis
 25997 [perl #43033] [PATCH] Silence warning
 13420 [R] Comparing MCMClogit, glm and BRUGS
 44333 Re: [svn:perl6-synopsis] r14401 - doc/trunk/de...
 62197 Re: [cc-community] Strayform - where artists a...
 21450 Re: svn commit: samba r23691 - in branches:\n\...

body

42791 One reason I rejoined this group was to try to...
 34315 it's also not unbiased.\n\n\n\n-----Original M...
 54248 Why not plot directly to a bitmapped format, u...
 35193 When Creative Commons published their Podcasti...
 37793 =====...
 32395 This thread prompts me to ask about something ...
 56651 On Wed, Jun 20, 2007 at 11:44:21AM -0500, Chri...
 30652 Author: allison\nDate: Sat Apr 14 17:06:35 200...
 49556 >From the help page\n\n 'pdf' writes unco...
 434 Thanx for your answer. I don't have the book, ...
 61015 Spie Sutherland wrote:\n> There seems to be se...
 53029 This will create a list of the matrix subsets:...
 1066 Hi Ben,\n\nYou can use shift for this.\n\nsub ...
 44225 -----BEGIN PGP SIGNED MESSAGE-----\nHash: SHA1...
 21039 Hi all,\nI have a little problem selecting som...
 62338 Send MSN-list mailing list submissions to\n\ntm...
 2380 On 4/11/07, Sundar Dorai-Raj wrote:\n\n> Seem...
 33977 Author: abartlet\nDate: 2007-07-04 03:25:44 +0...
 4286 what does sendmail have to do with the aliases...
 28368 Embarrassingly, I don't know awk or sed but R's...
 25997 # New Ticket Created by Steve Peters \n# Plea...
 13420 \nHello,\nI have two "related" questions, one ...
 44333 Darren Duncan wrote:\n> Jonathan Lang wrote:\n...
 62197 drew Roberts wrote:\n> On Friday 08 June 2007 ...
 21450 On Tue, Jul 03, 2007 at 11:34:02PM +0000, idra...

```
[262]: train[train_predicted == 'spam'][0:25]
```

```
[262]:
```

	to_address \
7375	theorize@plg.uwaterloo.ca
22780	theorize@plg.uwaterloo.ca
63457	"theorize" <theorize@plg.uwaterloo.ca>
26257	"Subscriber" <producttestpanel@speedy.uwaterlo...
58300	the00@plg2.math.uwaterloo.ca
41544	producttestpanel@flax9.uwaterloo.ca
1519	"Gnitpick" <gnitpick@flax9.uwaterloo.ca>
32192	<ktwarwic@flax9.uwaterloo.ca>
7250	<theorize@plg.uwaterloo.ca>
13006	<gnitpick@flax9.uwaterloo.ca>
8258	<manager@flax9.uwaterloo.ca>
42460	"The00" <the00@plg2.math.uwaterloo.ca>
1046	"Subscriber" <cruiseca@flax9.uwaterloo.ca>
5802	<mail@speedy.uwaterloo.ca>
59057	<smiles@speedy.uwaterloo.ca>
59295	1@bellsouth.net
24453	"Kami Crawford" <mail@speedy.uwaterloo.ca>
31894	theorize@plg.uwaterloo.ca
29012	mail@flax9.uwaterloo.ca
6031	mail@speedy.uwaterloo.ca
23167	ktwarwic@flax9.uwaterloo.ca
10491	<debian-legal@lists.debian.org>
6222	ktwarwic@flax9.uwaterloo.ca
1981	<the00@plg2.math.uwaterloo.ca>
493	gnitpick@flax9.uwaterloo.ca

	from_address \
7375	"Darla Blanco" <p4L4MLJ@fanciers.com>
22780	"Most languages" <stzezkh@ehotelier.com>
63457	"Fabian Jewell" <tiqannada@hosttune.com>
26257	"Shipment Notice #HL-225825" <Jenny@strongimpr...
58300	"Kimberley Mora" <the00@plg2.math.uwaterloo.ca>
41544	CollegePlanner <CollegePlanner@massiveservices...
1519	"Sydney Car Centre" <qaubm@altern.org>
32192	"HowStuffWorks" <newsletter@howstuffworks.com>
7250	"Guadalupe Rhodes" <ifuuow@boardiq.com>
13006	"Zachary Adams" <Tristan.Richardson@adultactio...
8258	"Alexander" <richard@prostateforum.biz>
42460	"SCC" <jjnvi@tenchiclub.com>
1046	"Healthcare Degrees Online" <qq5ci9zmi@daisyha...
5802	"Janine Lockett" <jpil@boston-ed.com>
59057	"Zachary Brown" <automacdev.com@lottozubotto.com>
59295	UK ONLINE PROMO <r_hoaglund@bellsouth.net>
24453	"Selene Harper" <macbrydejtfhi@163data.com.cn>

31894 fought by <elawsuit@cccfs.org>
 29012 Federico <ter53@icqmail.com>
 6031 "gambling" <pmqigufgumq@verizon.net>
 23167 "CNNMoney.com Alerts" <cnnalerts@mail.cnn.com>
 10491 <boletim@ecolatina.com.br>
 6222 "therein Dolan" <largesseLeach@ac9hold.com>
 1981 "Trevor Wood" <qsolo.com@usedcompanycars.com>
 493 "Roe, Darrin" <Scruggs9S@yourpositiveenergy.com>

subject \

7375 diagnoses malawi commonality__
 22780 Loan for a low month payment
 63457 you afton go imlaystown
 26257 Brand Samples shipping to your area =?UTF8?Q?=...
 58300 Domingo Buy your loved one Glashutte
 41544 adf, 10k scholarship giveaways - next drawing ...
 1519 Account managers vacant position in the Sydney...
 32192 HowStuffWorks Lifestyle Newsletter May 10, 2007
 7250 Global sale before new project opening!
 13006 Zachary Adams get it!
 8258 Need medicine? All here!
 42460 job - just for you. [letter id: KU32860911070A]
 1046 Find the right program for you
 5802 Janine - Viagra for you!
 59057 Lotto Tickets from 50 countries around the world
 59295 NOTIFICATION...MAIL
 24453 Be careful
 31894 Be no pest
 29012 look-out the price list
 6031 DIE RALLYE IS GESTARTET!
 23167 Lockheed employs D.C.'s last honest man
 10491 =?iso-8859-1?Q?=5BNS=5D_E-mail_Ecolatina_-_Tra...
 6222 Go out with a new babe tonight
 1981 Beware of fake pills
 493 solid stainless steel replicas PCXAW68

body

7375 Hello,\n\nVisit our new online store and save...
 22780 \n\n\n\n\n\nThank you for your loan request, w...
 63457 Does Size Matter'?\n____\n\n60% of WOMEN said...
 26257 ttp://www.lynxtrack.com/afclick.php?o=4343&b=d...
 58300 \nHighest qualities Replic Watches now !\n\n\n...
 41544 \n\n251839499 FAIR USE of Philadelphia Cit...
 1519 \n\n\n\n\nWhile we may have high ...
 32192 \nhowstuffworks® | lifestyle May 10, 2007\nRSS...
 7250 \n\n\n\n\n\n\nDear valued member.\n\n\nExtra d...
 13006 OEM software means no CD/DVD, no packing case,...

```

8258                                \n\n\n\n\n\n\n\n
42460 \n\n\n\nWhile we may have high expe...
1046  \n\n\n <!--U.S. military officials said they h...
5802  \n\n\n\n\n\n\nVIAGRA\nIf you have a problem ge...
59057 \n\n\n\n\n\n\nLotteryagent (TM) is the only on...
59295 \nREFERENCE NUMBER:UK/839030X2/14\nThis Email ...
24453 \n\n\n\n\n\n\n\n\n\nWe were cystic working num...
31894 \n\n\n\n\n\n\n\nlooking anxiously about her. Oh...
29012 \n\n\n\nhttp://carepicture.hk\n\n\nViagra\n10 ...
6031  \n\n\n\n\n\n\nAN ALLE FINANZINVESTOREN!\nDIESE A...
23167 \n\n\n\n\n\n\n\nThe Internet home of:\n\n\n\n\n...
10491 \n\n\n\nUntitled Document\n\n\n\n\n\n\n\n\n\n\n...
6222
1981  \nHi\n\n\n\n\n\n\nIm very sleepy, she blurtd ou...
493   \n\nRolex Replica order! Tracking #EI59 - 3481...

```

4. and
5. It looks like the R newsletter was more the in the ham category, and the college emails were much more spammy. We can tell these apart by the email addresses, containing 'uwaterloo' in very many of the spam emails.